

*Digital Publishing Tools
for Indexing, Text Mining, and Dataset Curation*

Overview

The rise of a digital publishing ecosystem has significantly expanded the possibilities for academic/scientific publishing, both in terms of the range of content provided (for instance, augmenting text documents with multimedia presentations) and the capabilities of reader experience in the context of text documents themselves (for instance, enhanced indexes, searching, and multi-document corpora). The PacTk toolkit provides software components and code libraries helping authors, researchers, and publishers to maximally leverage these new digital capabilities.

Use Case 1: Building Indexes to Enhance Search Capabilities

Ordinary text search simply tries to match one or more words within the exact text of a document. Curated indexes, by contrast, identify correlations between indexed terms and pages/paragraphs within the main text, even if the specific language in those paragraphs employs variant word-forms, or is conceptually relevant to the search term even within explicitly mentioning it. When authors or editors take the time to compile an index, it is wasteful to ignore those efforts when users seek to perform text searches while reading the main document.

To effectuate a more precise and conceptual search functionality, PacTk supports enhanced search capabilities that incorporate data from index files, when available. In this context “indexes” are not only human-readable pages at the back of a book, but also structured metadata that can be embedded in PDF files and deserialized by conformant PDF viewers (human-readable indexes may be automatically generated from such indexes). When someone reading a text document searches for a term similar to one from the index — or selects an index heading from a list derived via index metadata — the search algorithm can locate matches via this metadata instead of or in addition to performing a brute-force search within PDF character streams. Special-purpose GUI windows can be shown that allow users to fine-tune the search keywords and options (see Figure 1 for an example of how such dialog windows might appear).

PacTk similarly provides tools for compiling indexes. In addition to listing associations between index terms (also called “headings”) and page numbers or paragraph codes (also called “locators”), index metadata should model heading/subheading relations, “See also” cross-references, and “redirection” (“See”) cross-references. Dedicated GUI windows are necessary to allow authors and editors to examine index metadata conveniently, by navigating through headings in multiple ways (by order or creation, alphabetical sequence, or following parent/subheading and redirect connections) and juxtaposing index-entry windows with PDF page displays (see Figure 2 for one version of a single-entry dialog window).

Use Case 2: PDF Search for Creating and Updating Indexes

When curating an index, authors will often start with a list of terms which have clear relevance for a publication’s topic, and then search for instances of those headings in the document text. Although locators may often correspond to paragraphs without an exact match, a raw-text search can be a useful means to establish as provisional collection of locators that can be extended afterward.

When performing such initial matches, it is helpful to vary the order or precision of search terms. For example, a person’s name in an index will typically be listed and alphabetized via last name first, but will usually be matched in the text in the opposite order (or last name alone). Likewise, instances of a single word within a phrase may be sufficient to confirm that the relevant stretch of text is indeed relevant for the entire phrase as a concept (for example, “House” in lieu of “House of Representatives”). Given these points, PacTk provides a front-end to PDF search functionality which takes a provided search term (such as an index heading) and allows users to quickly rearrange or exclude words from the keyphrase to construct the specific match presented for PDF search. Successful matches are highlighted both in the PDF page view and within secondary windows that show the raw text located within the relevant PDF page (see Figure 3). Presenting this contextual information allows the person compiling an index to verify that text matches are indeed conceptual hits for the relevant index entry, as opposed to false positives.

The screenshot in Figure 3 actually represents a variation on index-building, where PacTk was utilized to migrate an index for the second edition of a previously-published manuscript. In this case, the first-edition index provided a provisional set of headings whose locators had to be updated with paragraph codes for the new book. We therefore employed a version of a PacTk GUI that had two pairs of page and raw-text views. In this variation, by matching terms against both documents, it is easy to confirm that a locator for the newer edition




Figure 1: An index entry dialog window (and context menu for single and meta-index views)

directly updates an index-entry from the prior book, by checking that the surrounded text is similar or identical. Once entries are updated between the two versions, then new headings (plus new locators for old terms) can be added.

Use Case 3: Customizing GUI Tools for Indexing

As shown in Figure 3, the GUI has two PDF view windows to show the two editions, plus a central window with a variety of entry lines, buttons, and check boxes, intended to automate the data transfer as much as possible. After identifying all of the index terms from the prior volume and the corresponding page references, the software is designed to work through index entries one at a time — either in alphabetized sequence or where users select a specific entry to start with. For each entry, users can click through the different page matches from the first edition, and the app searches for similar terms in the second edition. Above-mentioned search features allow users to fine-tune the search key to be more precise if needed — for instance, inverting a proper name (e.g., the entry "Durkheim, Emile" should be searched as "Emile Durkheim") or restricting the terms to the most important words (e.g., "'Diagnostic Impression' chart" can be reduced to "Diagnostic Impression"). Making these adjustments by clicking buttons is quicker and easier for the user than typing new search terms manually.

Once the search terms are matched against the new document, the code shows matches highlighted in both the first and second edition, and also shows the plain text of the corresponding files in separate windows. It's worth mentioning that the highlights are constructed via native GUI overlays, rather than employing the internal PDF highlight mechanism — the former is more flexible, and allows features like semi-transparency and color variation (which prove quite useful during a session where multiple entries and/or term-variants have been processed). With these features, the user is able to quickly check whether a match replicates and entry in the earlier index or else is a new locator. In either case the user can check a box, or alternatively a context menu option, which registers the match for the new document's




Figure 2: Dialog window for local and meta-index searches

index and tracks the correlation between the old and new index (the software automatically identifies correct paragraph codes if those codes, rather than page numbers, are used to generate an index). When ready, the user can confirm that the data for the current heading is complete and proceed to the next one. The software shown in Figure 3 then generates HTML code for that entry, which ultimately gets folded into the compiled second-edition index (which could be in formats such as Word, PDF, XML, etc.). The user has the option of automatically storing the entry data on a web back-end, or via FTP, which comes in handy if multiple people are working on an index at the same time.

Visible in Figures 2 and 3 are many GUI controls through which the user can work through a list of entries via buttons, check boxes, context menus, and “dock” widgets (that can be affixed to the main window or float as separate windows as desired). Some of these controls help to quickly process individual entries; others affect the search process by filtering the target search area (e.g., restricting to certain pages — such as a specific chapter — or just to footnotes, etc.). These various GUI features enable the user to process a single entry much more quickly, compared to the manual process. Often it is possible to confirm a match between the earlier and later document, and update the later index accordingly, in just a few seconds.

Whether building an index from scratch or updating an old manuscript, these GUI capabilities streamline the indexing process. Rather than typing index entries directly into an index file — manually adjusting details like italics, subentry lists, roman vs. arabic numerals, and so forth — the user could rely on mouse clicks and context menus (with minimal actual typing) to compile a data structure holding all index entries. A human-readable index is then generated from that structured data, or authors could submit the index as a data file instead, which could be incorporated into Meta-Indexes discussed below.

The PacTk GUI tools are designed in a flexible manner, with few external code dependencies and an emphasis on autonomous, modular design. That is, individual window components and/or algorithm libraries can be combined together and fine-tuned for specific book




Figure 3: GUI functionality for compiling and curating indexes

projects. In the example discussed here, our project involved migrating data from an older edition of a book to a new index. Other publications may have different needs — for example, syncing an index to a published data set; or compiling an aggregate index from a volume with individual chapter authors. Modular, native-compiled GUI controls allow new secondary windows to be implemented so as to support convenient user interaction, or automation, to complete tasks requisite for specific publication’s workflow (see the summary of BdQS below for more info on dynamic GUI customization).

Meta-Indexing and Cloud Services

Use Case 4: Annotated PDFs as SVG Files

PDF annotation capabilities (which include editing and comment boxes) are unfortunately very limited. Many PDF viewers cannot show annotations at all; those that do often present annotations in small, hard-to-read boxes. One solution to this problem is to convert PDF documents to a series of SVG files, that can be browsed as individual pages. SVG is a general-purpose format for creating 2D graphics, and SVG versions of PDF pages can be annotated with overlays that employ the full range of SVG capabilities, including all mathematically-describable paths and shapes; semi-transparent colors; scalable text; and mouse/keyboard user interactions. Figure 4 shows a feature where annotated PDFs are converted to SVG files so as to better visualize the annotations (which in ordinary PDF viewers tend to be reduced to poorly-formatted comment boxes). Here the comment-text is enlarged and color overlays clearly identify which sections of text correspond to specific annotations. Figure 5 reflects special programming for a project where annotations were grouped into different categories, modeling linguistic/discursive properties of the PDF-rendered content; a master index was then compiled linking specific pages to the annotation categories, with hyperref links to access the individual pages as desired.

One benefit of PDF-to-SVG conversion is that SVG files thereby become independent resources that can be individually hosted online. Most web browsers have PDF plugins, so that PDF files can be viewed without being downloaded, but there are no guarantees that users will also be able to see PDF annotations in the browser. Moreover, there is no way to construct a hyperref link to a single page, section, or chapter; only to the PDF as a whole. With SVG files, on the other hand, each page can have its own web URL, and links to single pages can be embedded in web pages (and also other PDF documents). Such web links can be useful for collaborative editing as well as referencing pages in files already published.

PacTk also supports technology for embedding web viewers in standalone native (i.e., desktop-style) applications. In this manner, authors/editors can preview the appearance and functionality of SVG annotations before they are disseminated online. Special-purpose native/JavaScript interop allows the host application to respond to a wider range of user actions than browser-based code in isolation, so PacTk tools can be used to modify and style annotations, edits, and comment boxes while working directly with SVG files.

No.	Field Name	Type	Description
30	NAICS 4	C	Fourth six-digit North American Standard Industry Classification System (NAICS) code entered by facility. NAICS codes reported by facilities from RY 2006 to present. NAICS codes in prior years were assigned by EPA. <i>Reference: Part I, Section 4.5b</i> <i>Maximum Length: 6</i>
31	NAICS 5	C	Fifth six-digit North American Standard Industry Classification System (NAICS) code entered by facility. NAICS codes reported by facilities from RY 2006 to present. NAICS codes in prior years were assigned by EPA. <i>Reference: Part I, Section 4.5b</i> <i>Maximum Length: 6</i>
32	NAICS 6	C	Sixth six-digit North American Standard Industry Classification System (NAICS) code entered by facility. NAICS codes reported by facilities from RY 2006 to present. NAICS codes in prior years were assigned by EPA. <i>Reference: Part I, Section 4.5b</i> <i>Maximum Length: 6</i>
33	DOCUMENT NUMBER	C	Number assigned to each TRI submission. <i>Reference: Part I, Section 4.5c</i> <i>Maximum Length: 13</i>
34	CHEMICAL	C	Name of the chemical or (generic name, if the chemical is claimed as a trade secret). <i>Reference: Part II, Section 1.2 or Part II, Section 1.3</i> <i>Maximum Length: 70</i>

Certain cells under this column contain the text "Nitrate compounds (water dissociable, reportable only when in aqueous solution)" without surrounding quotes, or employing a separator character other than the comma. As a result, these strings are parsed as two separate cells, causing all subsequent cells to the left, within the affected rows, to be shifted rightward; which of course corrupts any objects decoded from these rows.

Figure 4: Mapping PDF pages to SVG views for improved annotation support

Use Case 5: Curating Meta-Indexes

The value of indexes for enhancing local text search may certainly be extended to more general web/database searches.

Because PacTk represents indexes via metadata, it is straightforward to merge indexes from multiple documents into an aggregate database, or meta-index. Meta-indexes may be compiled for individual book series; academic journals; or even large-scale document corpora. Such meta-indexes could be made available to users in one of two ways. On the one hand, users might enter search terms and query against an overall bibliographic database, without necessarily starting from one specific book or publication. Alternatively, a user might be reading a specific digital publication and request a local search that then expands to include a meta-index. That is, users could indicate an interest in match-results not only for the book they are currently reading but also for pages in other volumes stored alongside it in a book series or corpus archive.

Hosting and curating a meta-index is facilitated by PacTk’s structured metadata format. Because all books in a series/archive will share a metadata profile, new indexes can be merged into a preexisting database as soon as they are finalized. PacTk provides an API protocol for registering a new index as well as querying for results in a hosted meta-index. PacTk provides native API tools so that PDF viewers can query against a database when expanding local to meta-index searches. Meta-index APIs could likewise be used for search front-ends




Figure 5: A category-based web-accessible index table for SVG pages

exposed to users via web pages for specific journals, book series, abstracts aggregators, and other publishers’ web content.

Existing bibliographic aggregators, such as Dimensions, pair API protocols with customized query languages for bibliographic databases. To support such technology, PacTk includes code for a Bibliographic-database Query Infrastructure (BdQS) for exchanging structured data with API hosts. Instead of REST-encoded URLs, more complex query prompts (and responses) can be encoded via BdQS for POST requests (native-compiled code can leverage easy-to-use request builder objects). The BdQS framework actually has several use cases within PacTk alongside API query factories; other use cases related to GUI programming were mentioned earlier.

The PacTk code base also includes prototype Cloud-Native service implementations for hosting meta-index databases. These cloud applications — which can be developed and tested as self-contained desktop components before being deployed online — encapsulate many details of meta-index management, database admin, and API support, with significant code reuse between native/desktop clients and API/cloud server endpoints.

Use Case 6: Machine-Readable Text Encoding

Because launching PDF text search for each API query would be very inefficient, cloud-based meta-indexes need to compile machine-readable representations of text document contents. This process also corrects for the inherent limitations of PDF text searches, which can be counter-productively stymied by such basic typographic conventions as ligatures, end-of-line hyphenation, apostrophes, footnote marks, and many other familiar presentational details.

With full metadata in place, PacTk indexes construct in-document locators that combine two different target origins: on the other hand, text spans are declared via machine-readable encoding for text mining, rather than (or in addition to) the internal PDF text representation, which is often an incomplete representation of the document text as understood by human authors and readers. At the same time, the visible outlines of text-segments that match index terms are stored via PDF page coordinates, which may be traced by SVG or other graphics overlays. In both cases, the internal PDF searching and highlighting functionality is supplanted by a two-facet text and graphics representation. Ordinary PDF search may be employed initially to “bootstrap” an index, but once finalized the document metadata can be consulted as a basis for searches both locally and via APIs.

Rigorous text encoding can also improve User Experience working locally with PDF files. For example, users could copy the text of a full sentence where a match occurs — not just the search term itself — with a single button-click. Figure 6 shows a range-select option




Figure 6: Flexible range-select features

that syncs mouse drags with key-presses to select a multi-line area that extends to the margins in the middle lines (a useful alternative to ordinary PDF ranges that can only be a simple rectangular box, thereby either cutting important words from left or right or including extraneous words around the range's start and end). Advanced GUI programming in conjunction with machine-readable text encoding allows users to accomplish their goals in find, copy, and remote-search functionality with less hassle than traditional PDF viewers (which rely on flawed internal text representation).

Use Case 7: Rigorous Text Mining

The limitations of PDF text encoding serve as a hindrance to advanced text-mining algorithms. These lacunae became evident during the Covid pandemic, when scientists sought to gather intelligence from a broad pre-pandemic literature about Coronaviruses and the first (2002-2004) SARS outbreak, as well as new publications which quickly emerged summarizing research in to SARS-CoV-2. The Allen Institute for AI (AI2), for example, compiled a large CORD-19 corpus using a project-specific JSON encoding of all document texts. However, AI2 pointed out that transcription errors were limiting the effectiveness of text-mining algorithms implemented against the corpus (an extended analysis of CORD-19 and other Covid-related data integration projects can be found in our book *Innovative Data Integration and Conceptual Space Modeling for COVID, Cancer, and Cardiac Care*, Elsevier 2022). We identified problems such as inconsistent representation of medical and chemical terms/formulas — such that correlations between publications were not properly identified — and ambiguities in modeling the structural breakdown of natural-language text in paragraphs and sentences. AI2 argued that new machine-readable text formats were necessary, rather than relying on PDF representations, and issued a “call for action” encouraging publishers and scientists to develop modernized text-encoding and metadata standards. We thereby developed a novel document-preparation system that generated structured text representations alongside conventional PDF files.

PacTk-enabled documents can utilize similar encoding to expose full text content, alongside index metadata, to bibliographic databases. Rigorous encoding allows databases to systematically address ambiguities related to match contexts and query results. For example, suppose an API request models a search for a keyphrase and requests all sentences where the phrase appears, among a multi-document collection. In order to fulfill the request, a query engine would need to clarify the proper representation of a “sentence,” as well as adjust for anomalies such as intraword hyphenation. For example, requests for complete sentences may or may not specify that the content of footnotes or block quotes inside the sentence should be included. Rhetorical features such as footnotes, quotes, figure illustrations, enumerations, and itemized lists all represent discursive structures that complicated the simplistic breakdown of text into sequences of paragraphs and sentences. Enhanced search, in particular, should allow queries to filter search targets into footnote text, section/chapter




Figure 7: Enhanced GUI capabilities enables by machine-readable text encoding

titles, quotes, content within figures/diagrams, and annotations/comments, as well as “main” document text. Query protocols should therefore recognize a plurality of discursive structures both to establish search parameters and to clarify the desired scope for contextual data. These details can be exposed to APIs via PacTk’s BdQS framework.

Data Publishing

Use Case 8: Deserializing Data Files

In contemporary publishing, text documents are often paired with scientific data sets in accord with FAIRsharing research-transparency principles (“FAIR” stands for Findable, Accessible, Interoperable, Reusable). For users intending to examine published data sets, raw data files are of limited value. Sometimes data sets can be loaded into proprietary software used by research labs, but many readers will not necessarily have access to such software (thereby failing the Accessibility and Reusability criteria). Consequently, specific scientific disciplines are often associated with domain-specific data publishing formats that are paired with dedicated deserialization tools (i.e., computer code which reads raw data files into live memory, without relying on special-purpose software). Such code libraries are often maintained by individual institutions or organizations, such as academic departments, on behalf of a broad research community. Examples include FASTQ and VCF for genomics; FITS and netCDF in astronomy and earth sciences; PDB and MOL in biochemistry; IFC in engineering; SEG-Y and BAG for geophysics and oceanography; CoNLL-U in computational linguistics; DICOM and OME-TIFF for bioimaging; and GIS-indexed attribute layers and shapefiles for geospatial data sets, among many others.

Some domain-specific file formats are derived from general-purpose formats such as CSV, XML, and JSON. Others use idiosyncratic text and/or binary serialization protocols, so that custom parsers are needed to consume byte or character streams from raw data files. Even when the actual encoding is performed via a generic standard like XML, dedicated deserialization code is still necessary — generic parsers can convert a given XML document to a Document Object Model (DOM) infoset, but algorithms need to navigate the document hierarchy and extract text strings from element nodes, which become the input for initializing object data fields.

In principle, supplying code libraries for parsing raw text files and/or traversing infosets embodies FAIR standards by ensuring that a typical reader who accesses published books or articles can likewise obtain usable research data. In practice, however, code libraries can lag behind the evolution of data standards, and FAIR access may be hindered by versioning conflicts, programming language barriers, operating-system conflicts, and similar obstacles. It is difficult to consistently maintain code libraries for a variety of different operating systems, coding languages, computational environments, and non-backwards-compatible versions of external dependencies, especially

when the data format targeted by a library itself goes through multiple changes.

PacTk seeks to address these concerns by engineering the BdQS framework to support data deserialization as well as bibliographic database queries. That is, deserialization algorithms can be expressed in BdQS and compiled to a custom Intermediate Representation/bytecode stack. BdQS is implemented in largely self-contained native code than can be dropped as source files into any **C++** project (or compiled into dynamic libraries invoked by scripting languages). As a result, BdQS can be ported to many diverse environments and embedded into many scientific applications, thereby ensuring that up-to-date deserialization capabilities are accessible for any file type for which a BdQS-based parser has been published. The entire BdQS stack could moreover be released as part of the source code contained in a published data set, so that deserialization functionality becomes part of the data set's overall package.

In general, BdQS deserializers can either parse character-streams directly — by registering callback functions responding to matches in Regular-Expression-like pattern language — or else post-process XML, JSON, CSV, and similar formats by declaring handlers for traversal events. BdQS provides a flexible suite of overload and redirection protocols to marshal post-parse navigation algorithms to **C++** methods that can granularly model preconditions for reading specific data fields (e.g., required nonempty, numeric range-validation, controlled vocabularies for string columns, etc).

Use Case 9: Data-Set Visualization and Nanopublications

Embedding deserialization code is a good illustration of FAIRsharing priorities: ideally, all the resources needed to (re)use a data set should be provided as part of the data package itself, rather than expecting users to download separate dependencies from external sources. Similar goals inform the design of data visualization features specific to individual data sets.

Some data sets might be loaded into readily-available software, but in many cases code for visualizing and studying raw data files should be provided as part of the data set itself. In these situations, the package can include source code files that, once compiled, produce an executable program which renders the concomitant data set in a visual, interactive manner. This may involve the use of charts or plots to show statistical trends as well as structured GUI controls, or separate windows, to show the attributes of individual data objects. For purposes of discussion, we refer to such self-contained resources specific to an individual data set as “Data Set Applications” (DSAs).

When a document (book or article) is explicitly paired with a published data set, information about the data set should be included in index/search metadata. This would allow keyword searches to be extended to raw data fields, as well as data-structure parameters (such as column labels, source-code annotations, procedure names, etc.). Moreover, sections of text can be isolated as cross-references for data-structure parameters as well as individual objects/records. For example, a specific column or field in tabular data structures can be cross-referenced with the paragraph in article text where that field is discussed (its observation/acquisition methodology, units of measurement, statistical distribution, etc.).

PacTk, in particular, envisions PDF documents links to supplemental resources that could take the form of graphics displays (showing things like statistical summaries) or dedicated GUI windows (showing, for example, structured data in tabular or hierarchical fashion). This model of supplemental presentations comports with the emerging idea of “nanopublications,” which are independently citeable works that are typically smaller in scale than, but may be paired with, a research article. Depending on one's preferred definition, an interactive statistical diagram, or a GUI window that can browse through discrete objects/records in a data set, may be considered a nanopublication or an integrated collection of multiple nanopublications. In either case, PacTk allows the nanopublications associated with a research document to be accessed interoperability from PDF (or SVG) files via the same conventions as indexing and text searches; and allows nanopublications to be integrated with meta-index databases for holistic corpora.

Use Case 10: Microcitations and Multimedia Cross-Referencing

Cross-referencing between research papers and data-set nanopublications is a special case of multi-media cross-referencing, where annotations defined in the context of a specific media type assert semantic or thematic relations to content with a different modality. Concrete examples of multimedia cross-referencing include geotagging video frames; linking CAD digital twins with product specs; annotating Regions of Interest within image series with coded labels linked to published raw data; embedding tags in 360° displays whose content includes URL strings that encode identifiers for data set objects; developing color- and/or shape-coded attribute layers for GIS overlays which allow users to browse between digital maps and structured object displays; and constructing numeric tuples from 3D graphics scenes such that particular camera orientation/location and zoom levels can be recorded as a unit, allowing the relevant 3D content to be repeatedly visualized from a specific perspective (in other words, links can be embedded in PDF files, web pages, software, or any other context so that users can reconstruct a specific 3D state).

The challenge when implementing multimedia cross-references is to ensure that viewers for two divergent media types can properly interoperate. With these cross-references, users should be able to navigate from a window dedicated to one media/file to a window rendering a different format, and correctly isolate the portion of a file (not just the overall file) embodied in a cross-reference. For example, geotagging video frames should allow users to pause a video at a particular location and then, via the geotag, switch to a map display

centered on the geospatial coordinates visible in the video (consider, e.g., a video documenting progress in an environmental remediation site). Conversely, GIS attribute layers can represent locations for which video content is available, and allow the users to watch the relevant videos starting from the appropriate time-point. Such interoperability is only possible if the components and/or applications responsible for rendering the two content-types share a common annotation and message-passing protocol.

The PacTk code base includes prototype viewers for some widely-used media types, including video (and audio) players, digital maps, 360°/panoramic photography, 3D triangular-mesh displays, image series, and parse-representation formats associated with text mining and computational linguistics. These prototypes illustrate how cross-referencing protocols may be adopted, and can be extended to provide domain-specific plugins for PDF viewers or bundled into data-set application code. More generally, PacTk defines a cross-referencing protocol that ensured proper interoperability between multimedia displays and text/bibliographic databases. This protocol can then be adapted for other media types, including special-purpose formats specific to individual scientific disciplines (such as 3D displays for molecular visualization, in cheminformatics; Digital Model Repositories for oncology and other biomedical fields; radiomic biomarkers for diagnostic imaging; and audio files with waveform visualizers for spoken-language corpora).

Standardized cross-referencing protocols can allow scientific applications to interoperate with the range of assets that may be published alongside research works, as supplemental materials. Rather than relying on generic PDF viewers, scientists could then employ PDF applications specifically engineered to work with the data formats and visualization strategies specific to a given research discipline.

Video Links

For more information, please see our software demo videos: